

Data in Motion – Intelligent Multi-Cloud Data Management

Project Presentation (Hackathon Prototype)

Author: Team GigaBITS

Date: 2025-11-08

Version: 0.1

1) Problem Understanding

Organizations operate across on-prem, private cloud, and public clouds. They need to:

- Optimize where data lives (hot/warm/cold) to balance performance and cost
- Migrate data across clouds with minimal disruption
- Stream, analyze, and act on real-time access patterns
- Predict usage trends and proactively move or reclassify data
- Maintain availability and consistency during failures
- Offer operators a unified view of data placement, costs, and migration activity

This prototype showcases an event-driven data management platform that uses access patterns and policies to automate placement decisions while ensuring integrity and availability.

2) Objectives

- Optimize Data Placement based on access frequency, latency SLOs, cost, and predicted trends
 - Enable Multi-Cloud Data Migration with integrity checks and minimal downtime
 - Integrate Real-Time Streaming (Kafka) for events and orchestration
 - Provide insights and automated recommendations based on recent access and policy constraints
 - Ensure Consistency & Availability across replicas
 - Ship a Unified Interface (REST + optional CLI) for observability and control
 - Containerized deployment for straightforward local and cloud simulation
-

3) Architecture Overview

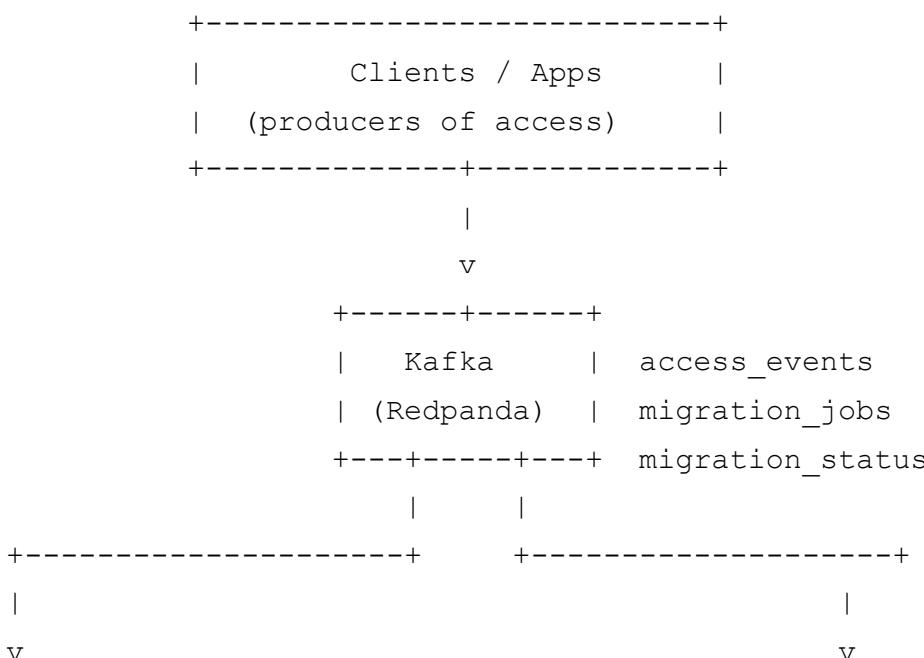
Core Components:

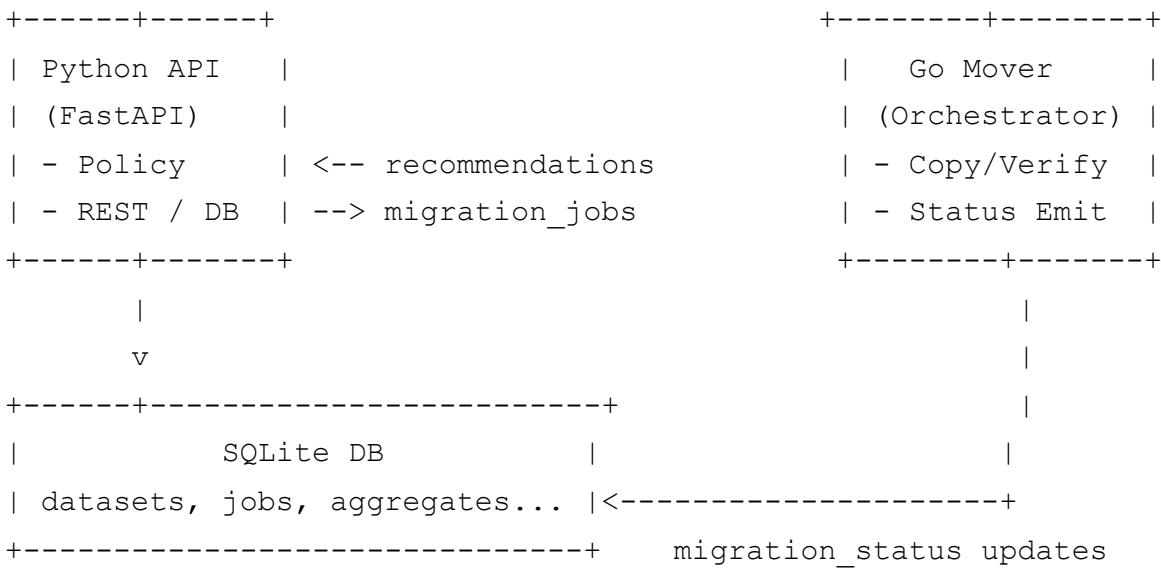
- Python FastAPI Service (Policy + Analytics)
 - Ingests access events
 - Aggregates usage and applies rules for tier recommendations
 - Enqueues migration jobs and tracks status
 - Exposes REST APIs for datasets, policies, recommendations, and jobs
- Go Migration Mover
 - Consumes migration jobs (Kafka)
 - Executes copy → verify → switch → complete cycle
 - Publishes migration status updates
 - Exposes REST APIs for job status and health
- SQLite Metadata Store
 - Tracks datasets, jobs, and aggregates
- Kafka-compatible Broker (Redpanda for local dev)
 - Topics: access_events, migration_jobs, migration_status, recommendations
- Storage Abstraction
 - Local FS(file://) for on-prem simulation
 - Optional GCS(gs://) and MinIO(s3://) integrations
- Observability
 - Health endpoints, structured logs, and metrics

High-Level Flow:

1. Clients send access events → Python API → Kafka (access_events)
 2. Python aggregates usage and applies rules → emits migration_jobs
 3. Go mover consumes jobs → runs migration lifecycle → emits migration_status
 4. Python updates DB → recommendations and dashboards reflect current state
-

4) Architecture Diagram (ASCII)





5) Components & Responsibilities

- Python FastAPI:
 - Endpoints: /datasets, /access-events, /recommendations, /plan-migration, /jobs, /train
 - Produces: access_events, migration_jobs
 - Consumes: migration_status (updates SQLite)
 - Tiering logic: recency-based heuristic for recommendations
- Go Mover:
 - Consumes: migration_jobs
 - Produces: migration_status
 - Lifecycles: queued → copying → verifying → switching → completed/failed
 - Persistent state: jobs mirrored in SQLite
- Kafka (Redpanda):
 - Simple single-node broker for dev
- SQLite:
 - Lightweight metadata store for rapid iteration

6) Data Model (SQLite, Minimal Viable)

Tables (current):

- datasets
 - id, name, path_uri, current_tier(hot|warm|cold), latency_slo_ms, size_bytes, owner, last_access_ts, created_at, updated_at
- migration_jobs
 - job_id(UUID), job_key(unique), dataset_id, source_uri, dest_uri, dest_storage_class
 - status(queued|copying|verifying|switching|completed|failed|enqueue_failed)
 - error, submitted_at, updated_at

Roadmap tables (extensible in Go storage.go + future ORM):

- replicas, policies, access_aggregates, recommendations, models
-

7) Streaming Topics

- access_events: dataset access telemetry (read/write, size, client latency, timestamp)
 - migration_jobs: commands for Go mover (source/dest, class)
 - migration_status: lifecycle updates for jobs (copying, verifying, switching, completed, failed)
 - recommendations (optional): human-readable recommendations from Python
 - metrics (optional): counters/latency from services
-

8) Placement Strategy

- Signals:
 - read_count, write_count, bytes_read, bytes_written
 - last_access_age_s, size_bytes
 - policy constraints (latency SLOs, cost targets)
 - Decision rules:
 - Recency threshold
 - <=60s → hot
 - <=10m → warm
 - 10m → cold
 - Policy checks may adjust the final tier
 - Metadata:
 - Decisions and reasons are stored in SQLite
 - Operation:
 - Recommendations can be computed on demand via the API and the Web UI
-

9) Migration Lifecycle

- queued
- copying: stream source → dest; compute checksum
- verifying: compare checksum/ETag; ensure object count
- switching: atomically update primary replica and dataset pointer
- completed: source retired if policy allows
- failed: supports retry; idempotent job_key prevents duplication

Consistency & Idempotency:

- Source remains primary until verification passes
 - Job keys provide dedupe; retries safe
 - Transactions on DB updates
-

10) Security & Compliance (MVP → Production)

- In-flight: TLS(future); local dev simple HTTP
 - At-rest: GCS/S3 native encryption; optional local envelope encryption
 - Access control: Add API keys or OAuth proxy in front of FastAPI/Go services
 - Auditing: Persist decisions(reason, model_version) in recommendations table
 - Secrets: Mount cloud credentials as read-only; never commit to repo
-

11) Cloud Simulation (GCP Free Tier)

- Use GCP Cloud Storage (Always Free) for “warm/cold” simulation
 - Local filesystem for “hot/on-prem”
 - Environment:
 - GOOGLE_APPLICATION_CREDENTIALS=/secrets/gcp-sa.json
 - Optional bucket: gs:// in eligible region
 - For local multi-cloud:
 - MinIO (S3-compatible) as private cloud simulation
-

12) Demo Plan (Local, Docker Compose)

1. Start stack:
 - docker compose up –build -d
 - Health: curl <http://localhost:8080/health> and <http://localhost:8090/health>
2. Create dataset:
 - POST /datasets with path_uri=file:///shared_storage/ds1
3. Generate access:
 - POST /access-events (multiple times with reads/writes)
4. Get recommendations:
 - GET /recommendations?dataset_id=1
5. Approve migration:

- POST /plan-migration (dest=file:///shared_storage/migrated/ds1)
- Observe Go mover transitions via GET /jobs and logs

6. Verify:

- Check dataset's current_tier and job completion status
-

13) Performance & Insights (Prototype-Level)

- SQLite with WAL is sufficient for serialized metadata ops in dev
 - Kafka-based eventing enables decoupled, asynchronous orchestration
 - Go mover can scale out horizontally with multiple consumers and partitioned topics
 - Tiering logic is heuristic and can evolve over time
-

14) Scalability & Roadmap

Near-term:

- Finish full DB wiring for all endpoints; persist recommendations
- Implement real Kafka consumer in Go (segmentio/kafka-go or sarama)
- File-based migration: chunked copy + checksum verification

Mid-term:

- Switch metadata to Postgres/Cloud SQL
- Add Prometheus metrics and Grafana dashboards
- Add policy-as-code and cost modeling engine
- Kubernetes deployment (HPA for mover replicas)

Long-term:

- Storage drivers for GCS/S3 with optimized transfers (multi-part, retries)
 - Advanced feature sets (seasonality, client clusters) and richer rule-based scoring
 - Multi-region replication with conflict resolution strategy
-

15) Risks & Mitigations

- SQLite contention → WAL + move to Postgres
 - Kafka outages → buffered retries; NullKafka fallback for dev
 - Partial migrations → idempotent job processing; resumable copy design
 - Credential security → secret mounts; least-privilege service accounts
 - Model drift → scheduled evaluation & retraining; fallback to heuristic
-

16) API Summary (Draft)

Python API:

- GET /health
- POST /datasets; GET /datasets; GET /datasets/{id}; PATCH /datasets/{id}
- POST /access-events
- GET /recommendations[?dataset_id=...]
- POST /plan-migration
- POST /train
- GET /jobs; GET /jobs/{job_id}

Go API:

- GET /health
 - GET /jobs; GET /jobs/{id}
 - POST /jobs/retry/{id}
-

17) Quickstart (Local)

- docker compose -f deploy/docker/docker-compose.yml up –build -d
 - Create dataset:
 - curl -X POST <http://localhost:8080/datasets> -H 'Content-Type: application/json' -d '{"name":"sample","path_uri":"file:///shared_storage/sample","size_bytes":1048576}'
 - Access event:
 - curl -X POST <http://localhost:8080/access-events> -H 'Content-Type: application/json' -d '{"dataset_id":1,"op":"read","size_bytes":4096,"client_lat_ms":12.3}'
 - Recommendation:
 - curl http://localhost:8080/recommendations?dataset_id=1
 - Plan migration:
 - curl -X POST <http://localhost:8080/plan-migration> -H 'Content-Type: application/json' -d '{"dataset_id":1,"target_location":"file:///shared_storage/migrated/sample","storage_class":"standard"}'
 - Observe jobs:
 - curl <http://localhost:8080/jobs>
 - curl <http://localhost:8090/jobs>
-

18) Closing

This prototype demonstrates an intelligent, event-driven foundation for multi-cloud data management:

- Streaming-first architecture with decoupled services
- Data-driven, policy-aware placement decisions

- Safe, idempotent migration lifecycle
- Extensible schema and storage drivers
- Containerized for rapid local and cloud simulations

We're ready to extend toward real cloud storage operations (GCS/S3) and robust observability to meet production needs.

Thank you!